

# Assingment 1

*by Ebba Bergman*

Let's do something very similar to the lab

**Hand in:** This notebook, and a pdf of this notebook. No written answers to the questions are required, they are only here to help you learn

**You are free to discuss the general concepts with other groups, but we encourage you not to exchange code for your own learning**

A lot of the code below is inspired labs developed by Christophe Avenel at NBIS , labs and assignments made by Phil Harrison as well as by <https://www.tensorflow.org/guide/keras/functional/>,

In [6]: *## First we need to import all of the packages we need*

```
import numpy as np
import tensorflow as tf
import pandas as pd
from PIL import Image
import IPython
from tensorflow import keras
from tensorflow.keras import layers
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix

from tensorflow.keras.preprocessing.image import ImageDataGenerator

import cnn_helper
```

Note: the cnn\_helper was written by Christophe Avenel, and his code (including his lab which this one is based on), is available here: [https://github.com/NBISweden/workshop-neural-nets-and-deep-learning/tree/master/session\\_convolutionalNeuralNetworks/Labs](https://github.com/NBISweden/workshop-neural-nets-and-deep-learning/tree/master/session_convolutionalNeuralNetworks/Labs)

```

In [7... def plot_history(model_history, model_name):
    fig = plt.figure(figsize=(15, 5), facecolor='w')
    ax = fig.add_subplot(131)
    ax.plot(model_history.history['loss'])
    ax.plot(model_history.history['val_loss'])
    ax.set(title=model_name + ': Model loss', ylabel='Loss', xlabel='')
    ax.legend(['train', 'valid'], loc='upper right')

    ax = fig.add_subplot(132)
    ax.plot(np.log(model_history.history['loss']))
    ax.plot(np.log(model_history.history['val_loss']))
    ax.set(title=model_name + ': Log model loss', ylabel='Log loss',
    ax.legend(['Train', 'Test'], loc='upper right')

    ax = fig.add_subplot(133)
    ax.plot(model_history.history['accuracy'])
    ax.plot(model_history.history['val_accuracy'])
    ax.set(title=model_name + ': Model accuracy', ylabel='Accuracy',
    ax.legend(['train', 'valid'], loc='upper right')
    plt.show()
    plt.close()

    plt.savefig("History Plot.png")

```

## Set up the data, look at it

```

In [8]: ## Set up where to find our data
    data_directory = "./LabData/bloodcells_small/data/"
    labels_path = "./LabData/bloodcells_small/labels.csv"


In... # This is a dataframe, a way to look at data as tables.
    #Google "Python pandas dataframe" to get more information, or to find n
    # Anything you can do with data frames you could do with loops, but it
    df_labels = pd.read_csv(labels_path)

```

Q: Look at the labels, what columns do you think contains the true label?

```
In [1...  ## Let's look at the images - always a good start to the project  
        # Here random images will be displayed, run this several time to see  
  
        figure, ax = plt.subplots(2, 3, figsize=(14, 10))  
        figure.suptitle("Examples of images", fontsize=20)  
        axes = ax.ravel()  
  
        df_images_to_show = df_labels.sample(8)  
  
        for i in range(len(axes)):  
            row = df_images_to_show.iloc[[i]]  
            random_image = Image.open(data_directory + row["Filenames"].valu  
            axes[i].set_title(row["Class"].values[0], fontsize=14)  
            axes[i].imshow(random_image)  
            axes[i].set_axis_off()  
  
        plt.subplots_adjust(wspace=0.05, hspace=0.05)  
        plt.show()  
        plt.close()
```

## Examples of images



Q: Can you see any difference between the classes?

Q: Do you think a human being able to see the difference between classes makes it an easier or more difficult problem for a neural network?

```
In [11]: # What's the shape of the image?
         image_shape = np.array(random_image).shape
         print(image_shape)
```

```
(32, 32)
```

```
In [12]: # Let's look a little bit into the labels
         set_size = df_labels.size
         print(set_size)
         print(df_labels.head())
```

41578

	Filenames	Class
0	CRF022_T_1_ch5_106.png	T
1	CRF022_T_1_ch5_119.png	T
2	CRF022_T_1_ch5_123.png	T
3	CRF022_T_1_ch5_128.png	T
4	CRF022_T_1_ch5_134.png	T

```
In [13]: df_labels['Class'].value_counts()
```

```
Out[13]: neutrophil    4500
         monocyte      4303
         T             4100
         B             4032
         eosinophil    3854
         Name: Class, dtype: int64
```

## Divide the data for training, validation and test

```
... ## Next, let's divide the filtered rows into a train, validation and a test set.
class_column_header = "Class"
df_to_use = df_labels.copy() #We're copying the df_labels so that you can use it again

test_set_fraction = 0.1
validation_set_fraction = 0.2

df_test = df_to_use.groupby(class_column_header).sample(frac = test_set_fraction)
df_to_use = pd.concat([df_to_use, df_test, df_test]).drop_duplicates(keep='first')
df_valid = df_to_use.groupby(class_column_header).sample(frac = validation_set_fraction)
df_train = pd.concat([df_to_use, df_valid, df_valid]).drop_duplicates(keep='first')
```

```
In [15]: print(df_test.head())
```

	Filenames	Class
19071	CRF132_B_1_ch5_73.png	B
17414	CRF034_B_2_ch5_55.png	B
19390	CRF132_B_3_ch5_15.png	B
18484	CRF074_B_3_ch5_23.png	B
17214	CRF022_B_3_ch5_30.png	B

```

In... ## Set up generators that specify how the images are loaded, how many at
      ## that the images should be shuffled etc.
      batch_size = 8

      filename_column = 'FileNames'
      true_value = "Class"
      # create a data generator

      ## Note: we tend to get better results if the values of the pixels are b
      train_data_generator = keras.preprocessing.image.ImageDataGenerator(resc
      valid_data_generator = keras.preprocessing.image.ImageDataGenerator(resc
      test_data_generator = keras.preprocessing.image.ImageDataGenerator(resca

      train_generator = train_data_generator.flow_from_dataframe(
          df_train, directory=data_directory, x_col=filename_column, y_col=tru
          weight_col=None, class_mode='categorical', batch_size=batch_size, ta
      )

      valid_generator = valid_data_generator.flow_from_dataframe(
          df_valid, directory=data_directory, x_col=filename_column, y_col=tru
          weight_col=None, class_mode='categorical', batch_size=batch_size, ta
      )

      test_generator = test_data_generator.flow_from_dataframe(
          df_test, directory=data_directory, x_col=filename_column, y_col=true
          weight_col=None, class_mode='categorical', batch_size=batch_size, tar
      )

      train_steps=train_generator.n//train_generator.batch_size if train_gener
      validation_steps=valid_generator.n//valid_generator.batch_size if valid_

      Found 14968 validated image filenames belonging to 5 classes.
      Found 3743 validated image filenames belonging to 5 classes.
      Found 2078 validated image filenames belonging to 5 classes.

```

## CNN

Convolutional Neural Networks revolutionized the field of deep learning. You have seen how convolutions work in the lectures. One of the huge benefits of convolutions is that as the filters (sometimes called kernels in codes) move across the image the position of an object in an image becomes much less important than when we flattened images to use in traditional Artificial Neural Networks.

For this part of the lab you will try a couple of different architectures and hyperparameters. The **architecture** is basically the structure of the network: how many nodes, how many layers, and overall shape of these. The **hyperparamters** are most easily defined as all of the parameters changed *before* the training of the network begin, such as the number of epochs, what activation function to use in each layer, and which optimization method we use for backpropagation.

```
In [1... ## Set up the model architecture
# See https://www.tensorflow.org/guide/keras/functional/ if you want

cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(1, kernel_size=(3, 3), strides=1,padding='same')(cnn_inputs)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)
```

```
In [2... ## Define the model as a keras model
cnn_model = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs, name='cnn_model')
```

```
In [... ## We'll use the same generators as above here, so no need to redefine
## compile model
```

```
cnn_model.compile(optimizer=keras.optimizers.Adam(), loss='categorical_crossentropy')
cnn_model.summary()
```

Model: "cnn\_Model\_1"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d (Conv2D)	(None, 32, 32, 1)	10
max_pooling2d (MaxPooling2D)	(None, 16, 16, 1)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 5)	1285
Total params: 1,295		
Trainable params: 1,295		
Non-trainable params: 0		

```
In [26]: ## Actually train model
epochs = 8
history = cnn_model.fit_generator(generator=train_generator,
                                steps_per_epoch= train_steps,
                                validation_data= valid_generator,
                                validation_steps= validation_steps,
                                epochs= epochs)
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

Train for 1871 steps, validate for 467 steps

Epoch 1/8

1871/1871 [=====] - 20s 11ms/step - loss: 0.5173

- accuracy: 0.8115 - val\_loss: 0.5448 - val\_accuracy: 0.7875

Epoch 2/8

1871/1871 [=====] - 14s 8ms/step - loss: 0.5062 -

accuracy: 0.8140 - val\_loss: 0.5234 - val\_accuracy: 0.8116

Epoch 3/8

1871/1871 [=====] - 15s 8ms/step - loss: 0.4984 -

accuracy: 0.8167 - val\_loss: 0.5334 - val\_accuracy: 0.8001

Epoch 4/8

1871/1871 [=====] - 15s 8ms/step - loss: 0.4922 -

accuracy: 0.8183 - val\_loss: 0.5192 - val\_accuracy: 0.8086

Epoch 5/8

1871/1871 [=====] - 32s 17ms/step - loss: 0.4860

- accuracy: 0.8226 - val\_loss: 0.5196 - val\_accuracy: 0.8065

Epoch 6/8

1871/1871 [=====] - 14s 7ms/step - loss: 0.4832 -

accuracy: 0.8265 - val\_loss: 0.5107 - val\_accuracy: 0.8065

Epoch 7/8

1871/1871 [=====] - 14s 7ms/step - loss: 0.4779 -

accuracy: 0.8283 - val\_loss: 0.5157 - val\_accuracy: 0.8118

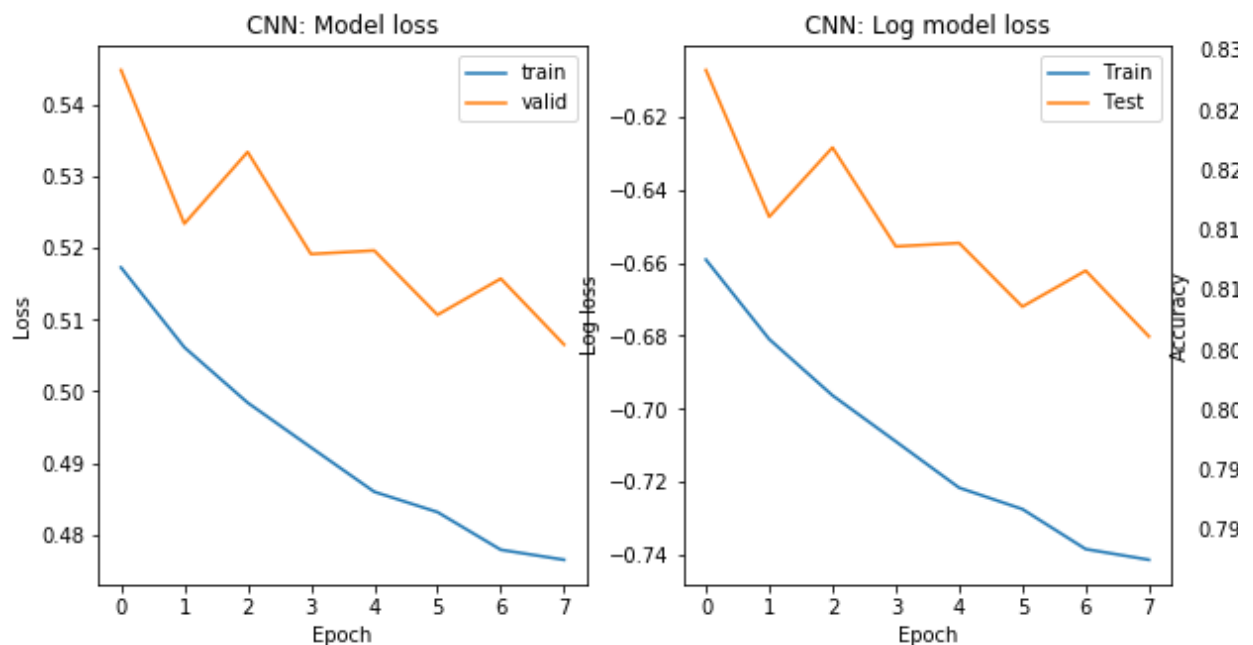
Epoch 8/8

1871/1871 [=====] - 15s 8ms/step - loss: 0.4765 -

accuracy: 0.8264 - val\_loss: 0.5065 - val\_accuracy: 0.8191

In [27]: *## Plot results*

```
plot_history(history, "CNN")
```




<Figure size 432x288 with 0 Axes>

In [28... *# plot confusion matrix*

```
cnn_helper.plot_confusion_matrix_from_generator(cnn_model, valid_ge
```



Accuracy: 0.8185947101255677



Q: What do the curves tell you about the models?

You can see some examples of how curves can look at : [https://uppsala.instructure.com/courses/ 2 3 8 0 4 /pages/deep-learning-plots/edit](https://uppsala.instructure.com/courses/23804/pages/deep-learning-plots/edit)

## Expanding the models

### Deeper models

Sometimes a deeper model and/or a more complex model, can be helpful. Try adding some more convolution layers and pooling layers to the model. Try changing the filter sizes, and the number of filters as well. More information about the convolutional layer can be found here:

[https://keras.io/api/layers/convolution\\_layers/convolution 2 d/](https://keras.io/api/layers/convolution_layers/convolution_2d/), maxpooling here: [https://keras.io/api/layers/pooling\\_layers/max\\_pooling 2 d/](https://keras.io/api/layers/pooling_layers/max_pooling_2d/), and a different kind of way of making models can be found here: <https://www.tensorflow.org/tutorials/images/cnn> and here <https://www.tensorflow.org/tutorials/quickstart/advanced>

```

In [... ] ## Set up the model architecture
## v.1.1
# Accuracy validation: 0.9419
# Comment: Could run more than 10 epochs.
##

#Change the code below so that the new model has roughly the same num
# Hint: you can add both more Conc2D layers, and increase the kernel

cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same', ac
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

## Define the model
cnn_model = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs, name=

## Compile the model
cnn_model.compile(optimizer=keras.optimizers.Adam(), loss='categorical_crossentropy')
cnn_model.summary()

```

Model: "cnn\_Model\_2"

Layer (type)	Output Shape	Param #
=====		
input_50 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_78 (Conv2D)	(None, 32, 32, 5)	50
max_pooling2d_74 (MaxPooling)	(None, 16, 16, 5)	0
conv2d_79 (Conv2D)	(None, 16, 16, 5)	230
max_pooling2d_75 (MaxPooling)	(None, 8, 8, 5)	0
flatten_35 (Flatten)	(None, 320)	0
dense_38 (Dense)	(None, 5)	1605
=====		
Total params: 1,885		
Trainable params: 1,885		
Non-trainable params: 0		

```

In [ ...  ## Set up the model architecture
          # Larger kernel from 3 to 5
          ##
          # Accuracy validation: 0.9430
          # Comment: Hard to train for validationi the accuracy goes up and down
          ##

          #Change the code below so that the new model has roughly the same number of parameters
          # Hint: you can add both more Conv2D layers, and increase the kernel size

          cnn_inputs = keras.Input(shape=(32,32,1))
          x = layers.Conv2D(5, kernel_size=(5, 5), strides=1,padding='same', activation='relu')(x)
          x = layers.MaxPooling2D(pool_size=(2, 2))(x)
          x = layers.Conv2D(5, kernel_size=(5, 5), strides=1,padding='same')(x)
          x = layers.MaxPooling2D(pool_size=(2, 2))(x)
          x = layers.Flatten()(x)
          cnn_outputs = layers.Dense(5, activation='softmax')(x)

          ## Define the model
          cnn_model = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs, name='cnn_Model_2')

          ## Compile the model
          cnn_model.compile(optimizer=keras.optimizers.Adam(), loss='categorical_crossentropy', metrics=['accuracy'])
          cnn_model.summary()

```

Model: "cnn\_Model\_2"

Layer (type)	Output Shape	Param #
=====		
input_49 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_76 (Conv2D)	(None, 32, 32, 5)	130
max_pooling2d_72 (MaxPooling)	(None, 16, 16, 5)	0
conv2d_77 (Conv2D)	(None, 16, 16, 5)	630
max_pooling2d_73 (MaxPooling)	(None, 8, 8, 5)	0
flatten_34 (Flatten)	(None, 320)	0
dense_37 (Dense)	(None, 5)	1605
=====		
Total params: 2,365		
Trainable params: 2,365		
Non-trainable params: 0		
=====		

```

In [... ## Set up the model architecture
# Small kernel from 3 to 1
##
# Accuracy validation: 0.9349
# Comment: Nice
##

#Change the code below so that the new model has roughly the same num
# Hint: you can add both more Conv2D layers, and increase the kernel

cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(1, 1), strides=1,padding='same', ac
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Conv2D(5, kernel_size=(1, 1), strides=1,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

## Define the model
cnn_model = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs, name=

## Compile the model
cnn_model.compile(optimizer=keras.optimizers.Adam(), loss='categorical_crossentropy')
cnn_model.summary()

```

Model: "cnn\_Model\_2"

Layer (type)	Output Shape	Param #
=====		
input_43 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_64 (Conv2D)	(None, 32, 32, 5)	10
max_pooling2d_60 (MaxPooling)	(None, 16, 16, 5)	0
conv2d_65 (Conv2D)	(None, 16, 16, 5)	30
max_pooling2d_61 (MaxPooling)	(None, 8, 8, 5)	0
flatten_28 (Flatten)	(None, 320)	0
dense_31 (Dense)	(None, 5)	1605
=====		
Total params: 1,645		
Trainable params: 1,645		
Non-trainable params: 0		
=====		

```

In [...] ## Set up the model architecture
# Big pool size from 2 and 2 to 8 and 4. Results in a feature maps of
##
# Accuracy validation: 0.8454
# Increased epoch to 20, still got questionable results :/
##

#Change the code below so that the new model has roughly the same num
# Hint: you can add both more Conc2D layers, and increase the kernel

cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same', ac
x = layers.MaxPooling2D(pool_size=(8, 8))(x)
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(4, 4))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

## Define the model
cnn_model = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs, name=

## Compile the model
cnn_model.compile(optimizer=keras.optimizers.Adam(), loss='categorica
cnn_model.summary()

```

Model: "cnn\_Model\_2"

Layer (type)	Output Shape	Param #
=====		
input_41 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_60 (Conv2D)	(None, 32, 32, 5)	50
max_pooling2d_56 (MaxPooling)	(None, 4, 4, 5)	0
conv2d_61 (Conv2D)	(None, 4, 4, 5)	230
max_pooling2d_57 (MaxPooling)	(None, 1, 1, 5)	0
flatten_26 (Flatten)	(None, 5)	0
dense_29 (Dense)	(None, 5)	30
=====		
Total params: 310		
Trainable params: 310		
Non-trainable params: 0		

```

In [153]: ## Actually train model
          epochs = 10
          history = cnn_model.fit_generator(generator=train_generator,
                                           steps_per_epoch= train_steps,
                                           validation_data= valid_generator,
                                           validation_steps= validation_steps,
                                           epochs= epochs
                                           )

```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to  
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to  
['...']
```

Train for 1871 steps, validate for 467 steps

Epoch 1/10

1871/1871 [=====] - 20s 11ms/step - loss: 0.5338  
- accuracy: 0.8208 - val\_loss: 0.3285 - val\_accuracy: 0.8908

Epoch 2/10

1871/1871 [=====] - 14s 7ms/step - loss: 0.3059 -  
accuracy: 0.9020 - val\_loss: 0.2735 - val\_accuracy: 0.9111

Epoch 3/10

1871/1871 [=====] - 14s 8ms/step - loss: 0.2669 -  
accuracy: 0.9174 - val\_loss: 0.2565 - val\_accuracy: 0.9157

Epoch 4/10

1871/1871 [=====] - 14s 8ms/step - loss: 0.2486 -  
accuracy: 0.9224 - val\_loss: 0.2503 - val\_accuracy: 0.9208

Epoch 5/10

1871/1871 [=====] - 16s 8ms/step - loss: 0.2375 -  
accuracy: 0.9276 - val\_loss: 0.2319 - val\_accuracy: 0.9264

Epoch 6/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.2288 -  
accuracy: 0.9290 - val\_loss: 0.2299 - val\_accuracy: 0.9272

Epoch 7/10

1871/1871 [=====] - 14s 8ms/step - loss: 0.2221 -  
accuracy: 0.9323 - val\_loss: 0.2266 - val\_accuracy: 0.9267

Epoch 8/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.2166 -  
accuracy: 0.9334 - val\_loss: 0.2188 - val\_accuracy: 0.9299

Epoch 9/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.2125 -  
accuracy: 0.9339 - val\_loss: 0.2264 - val\_accuracy: 0.9277

Epoch 10/10


1871/1871 [=====] - 16s 8ms/step - loss: 0.2105 -  
accuracy: 0.9349 - val\_loss: 0.2187 - val\_accuracy: 0.9350

In [154... *## Plot results*


```
plot_history(history, "cnn_model")
```

```
# plot confusion matrix
```

```
cnn_helper.plot_confusion_matrix_from_generator(cnn_model, valid_ge
```



Accuracy: 0.9345444830349987



Try a couple of deeper models and save your best one for further study

Add all these models beneath this heading

## Data Augmentation

Let's try something else, maybe you would like to add some data augmentation? Data augmentation basically means that we randomly alter the incoming images in different ways to make sure that the network can handle those types of variations.

If you want to read more you can look at this article, especially the "Data Augmentations based on basic image manipulations Geometric transformations" is of interest here: <https://journalofbigdata.springeropen.com/articles/10.1186/s40537-019-0197-0>

See [https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image/ImageDataGenerator](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator) for things you can try by adding input parameters to the `ImageDataGenerator()`.

Update the cell below to **include data augmentations, only in the training data generator then run your CNN again**



```

In... ## Set up generators
batch_size = 8

filename_column = 'FileNames'
true_value = "Class"
# create a data generator

## Note: we tend to get better results if the values of the pixels are b
train_data_generator = keras.preprocessing.image.ImageDataGenerator(resc
valid_data_generator = keras.preprocessing.image.ImageDataGenerator(resc
test_data_generator = keras.preprocessing.image.ImageDataGenerator(resca

train_generator = train_data_generator.flow_from_dataframe(
    df_train, directory=data_directory, x_col=filename_column, y_col=true
    weight_col=None, class_mode='categorical', batch_size=batch_size, ta
)

valid_generator = valid_data_generator.flow_from_dataframe(
    df_valid, directory=data_directory, x_col=filename_column, y_col=true
    weight_col=None, class_mode='categorical', batch_size=batch_size, ta
)

test_generator = test_data_generator.flow_from_dataframe(
    df_test, directory=data_directory, x_col=filename_column, y_col=true
    weight_col=None, class_mode='categorical', batch_size=batch_size, tar
)

train_steps=train_generator.n//train_generator.batch_size if train_gener
validation_steps=valid_generator.n//valid_generator.batch_size if valid_

Found 14968 validated image filenames belonging to 5 classes.
Found 3743 validated image filenames belonging to 5 classes.
Found 2078 validated image filenames belonging to 5 classes.

```

```

In [1... ## Set up the model architecture
### use your best model from above, and rename it here to cnn_model_e
cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same', ac
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

```

```

In [1... ## Define the model
cnn_model = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs, name

```

```

In [... ## Compile the model
cnn_model.compile(optimizer=keras.optimizers.Adam(), loss='categorica
cnn_model.summary()

```

Model: "cnn\_Model\_augmented"

Layer (type)	Output Shape	Param #
input_44 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_66 (Conv2D)	(None, 32, 32, 5)	50
max_pooling2d_62 (MaxPooling)	(None, 16, 16, 5)	0
conv2d_67 (Conv2D)	(None, 16, 16, 5)	230
max_pooling2d_63 (MaxPooling)	(None, 8, 8, 5)	0
flatten_29 (Flatten)	(None, 320)	0
dense_32 (Dense)	(None, 5)	1605
Total params: 1,885		
Trainable params: 1,885		
Non-trainable params: 0		

```
In [159]: ## Actually train model
          epochs = 10
          history = cnn_model.fit_generator(generator=train_generator,
                                           steps_per_epoch= train_steps,
                                           validation_data= valid_generator,
                                           validation_steps= validation_steps,
                                           epochs= epochs
                                           )
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

Train for 1871 steps, validate for 467 steps

Epoch 1/10

1871/1871 [=====] - 21s 11ms/step - loss: 0.5312

- accuracy: 0.8088 - val\_loss: 0.3096 - val\_accuracy: 0.8967

Epoch 2/10

1871/1871 [=====] - 21s 11ms/step - loss: 0.2787

- accuracy: 0.9088 - val\_loss: 0.2385 - val\_accuracy: 0.9237

Epoch 3/10

1871/1871 [=====] - 14s 7ms/step - loss: 0.2360 -

accuracy: 0.9244 - val\_loss: 0.2273 - val\_accuracy: 0.9275

Epoch 4/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.2204 -

accuracy: 0.9315 - val\_loss: 0.2362 - val\_accuracy: 0.9267

Epoch 5/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.2113 -

accuracy: 0.9317 - val\_loss: 0.1950 - val\_accuracy: 0.9403

Epoch 6/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.2015 -

accuracy: 0.9363 - val\_loss: 0.2018 - val\_accuracy: 0.9320

Epoch 7/10

1871/1871 [=====] - 14s 8ms/step - loss: 0.1992 -

accuracy: 0.9369 - val\_loss: 0.2450 - val\_accuracy: 0.9210

Epoch 8/10

1871/1871 [=====] - 14s 8ms/step - loss: 0.1937 -

accuracy: 0.9378 - val\_loss: 0.1964 - val\_accuracy: 0.9352

Epoch 9/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.1943 -

accuracy: 0.9391 - val\_loss: 0.1976 - val\_accuracy: 0.9371

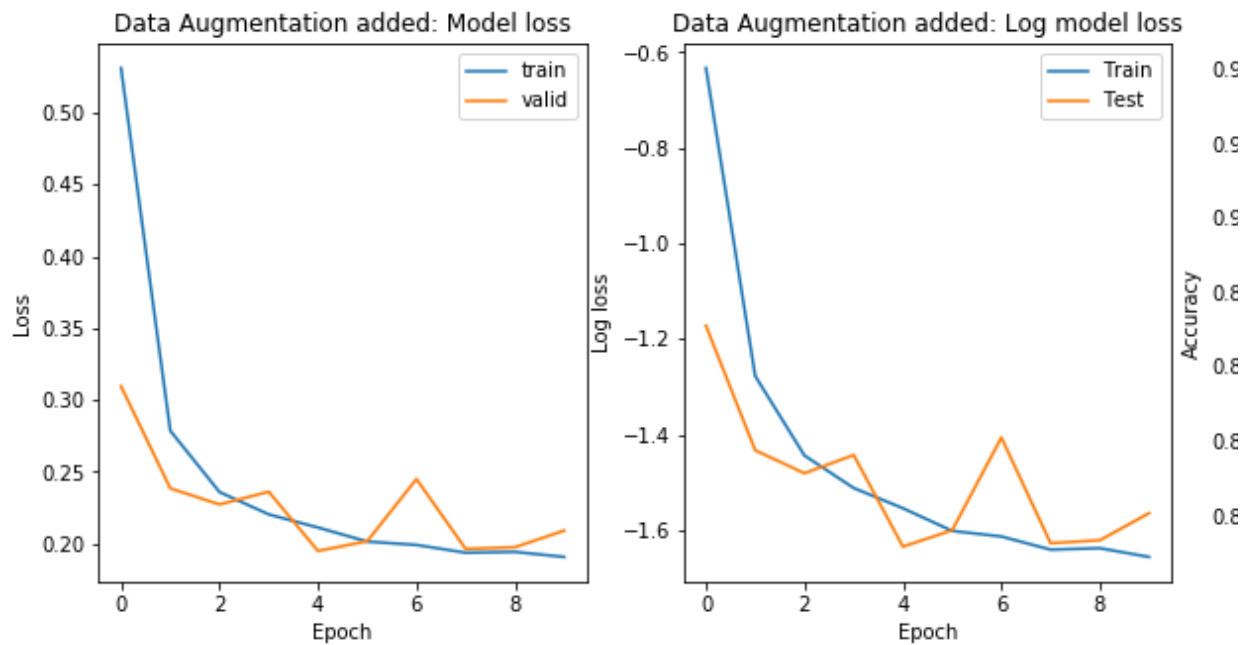
Epoch 10/10

1871/1871 [=====] - 16s 8ms/step - loss: 0.1908 -

accuracy: 0.9385 - val\_loss: 0.2091 - val\_accuracy: 0.9320

In [160]: *## Plot results*


plot\_history(history, "Data Augmentation added")



<Figure size 432x288 with 0 Axes>

```
In [161]: # plot confusion matrix
          cnn_helper.plot_confusion_matrix_from_generator(cnn_model, valid_gen)
```

Accuracy: 0.9318728292813251



Q: Did the data augmentation help? Why or why not? What makes this dataset more or less likely to be helped by data augmentation?

**Optional hints for question above**

1. Are the blood cells at random places in the image? No 2. Look at some of the images. Are the bloodcells centered? What could rotations or zooms change about this? It would change nothing 3. Are there color changes you could compensate for?

# Regularisation methods

Both BatchNormalization and DropOut are two different regularisation methods. Try adding both to the best working CNN model.

Read more about BatchNormalization here: [https://keras.io/api/layers/normalization\\_layers/batch\\_normalization/](https://keras.io/api/layers/normalization_layers/batch_normalization/) Read more about DropOut here: [https://keras.io/api/layers/regularization\\_layers/dropout/](https://keras.io/api/layers/regularization_layers/dropout/)

**Q: What are the main similarities and differences between these methods?**

```
In [1... # Create the model here
        ## Set up the model architecture
        ### use your best model from above

cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same', ac
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Dropout(.1)(x)
x = layers.BatchNormalization()(x)
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

In [1... cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same', ac
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Dropout(.2)(x)
x = layers.BatchNormalization()(x)
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

In [16... ## Define the model
cnn_model = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs, nam

In ... ## Compile the model

cnn_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0001
cnn_model.summary()
print(cnn_model.layers)
```

Model: "cnn\_Model"

Layer (type)	Output Shape	Param #
input_46 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_70 (Conv2D)	(None, 32, 32, 5)	50
max_pooling2d_66 (MaxPooling)	(None, 16, 16, 5)	0
dropout_8 (Dropout)	(None, 16, 16, 5)	0
batch_normalization_15 (Batch Normalization)	(None, 16, 16, 5)	20
conv2d_71 (Conv2D)	(None, 16, 16, 5)	230
max_pooling2d_67 (MaxPooling)	(None, 8, 8, 5)	0
flatten_31 (Flatten)	(None, 320)	0
dense_34 (Dense)	(None, 5)	1605

Total params: 1,905

Trainable params: 1,895

Non-trainable params: 10

```
[<tensorflow.python.keras.engine.input_layer.InputLayer object at 0x7f40a427fb70>, <tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f40a427f3c8>, <tensorflow.python.keras.layers.pooling.MaxPooling2D object at 0x7f40a427fa58>, <tensorflow.python.keras.layers.core.Dropout object at 0x7f40a43983c8>, <tensorflow.python.keras.layers.normalization_v2.BatchNormalization object at 0x7f40a4398828>, <tensorflow.python.keras.layers.convolutional.Conv2D object at 0x7f40a4398358>, <tensorflow.python.keras.layers.pooling.MaxPooling2D object at 0x7f40a46f60f0>, <tensorflow.python.keras.layers.core.Flatten object at 0x7f407c793d30>, <tensorflow.python.keras.layers.core.Dense object at 0x7f40540d1240>]
```

In [166]: *## Actually train model*

epochs = 15

```
history = cnn_model.fit_generator(generator=train_generator,
                                steps_per_epoch= train_steps,
                                validation_data= valid_generator,
                                validation_steps= validation_steps,
                                epochs= epochs
                                )
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

Train for 1871 steps, validate for 467 steps

Epoch 1/15

1871/1871 [=====] - 26s 14ms/step - loss: 1.4251

- accuracy: 0.3952 - val\_loss: 1.0239 - val\_accuracy: 0.6089

Epoch 2/15

1871/1871 [=====] - 15s 8ms/step - loss: 0.9046 -

accuracy: 0.6472 - val\_loss: 0.7296 - val\_accuracy: 0.7184

Epoch 3/15

1871/1871 [=====] - 15s 8ms/step - loss: 0.6930 -

accuracy: 0.7455 - val\_loss: 0.5983 - val\_accuracy: 0.7749

Epoch 4/15

1871/1871 [=====] - 15s 8ms/step - loss: 0.5814 -

accuracy: 0.7925 - val\_loss: 0.5229 - val\_accuracy: 0.8126

Epoch 5/15

1871/1871 [=====] - 16s 8ms/step - loss: 0.5226 -

accuracy: 0.8210 - val\_loss: 0.4852 - val\_accuracy: 0.8231

Epoch 6/15

1871/1871 [=====] - 15s 8ms/step - loss: 0.4790 -

accuracy: 0.8383 - val\_loss: 0.4559 - val\_accuracy: 0.8346

Epoch 7/15

1871/1871 [=====] - 14s 7ms/step - loss: 0.4555 -

accuracy: 0.8447 - val\_loss: 0.4131 - val\_accuracy: 0.8552

Epoch 8/15

1871/1871 [=====] - 34s 18ms/step - loss: 0.4238

- accuracy: 0.8570 - val\_loss: 0.3950 - val\_accuracy: 0.8603

Epoch 9/15

1871/1871 [=====] - 15s 8ms/step - loss: 0.4031 -

accuracy: 0.8646 - val\_loss: 0.3830 - val\_accuracy: 0.8643

Epoch 10/15

1871/1871 [=====] - 14s 8ms/step - loss: 0.3874 -

accuracy: 0.8709 - val\_loss: 0.3629 - val\_accuracy: 0.8721

Epoch 11/15

1871/1871 [=====] - 14s 8ms/step - loss: 0.3736 -

accuracy: 0.8777 - val\_loss: 0.3358 - val\_accuracy: 0.8838

Epoch 12/15

1871/1871 [=====] - 14s 8ms/step - loss: 0.3545 -

accuracy: 0.8815 - val\_loss: 0.3326 - val\_accuracy: 0.8844

Epoch 13/15

1871/1871 [=====] - 16s 8ms/step - loss: 0.3443 -

accuracy: 0.8860 - val\_loss: 0.3079 - val\_accuracy: 0.8964

Epoch 14/15

1871/1871 [=====] - 16s 8ms/step - loss: 0.3353 -

accuracy: 0.8912 - val\_loss: 0.2894 - val\_accuracy: 0.9028

Epoch 15/15

1871/1871 [=====] - 15s 8ms/step - loss: 0.3208 -

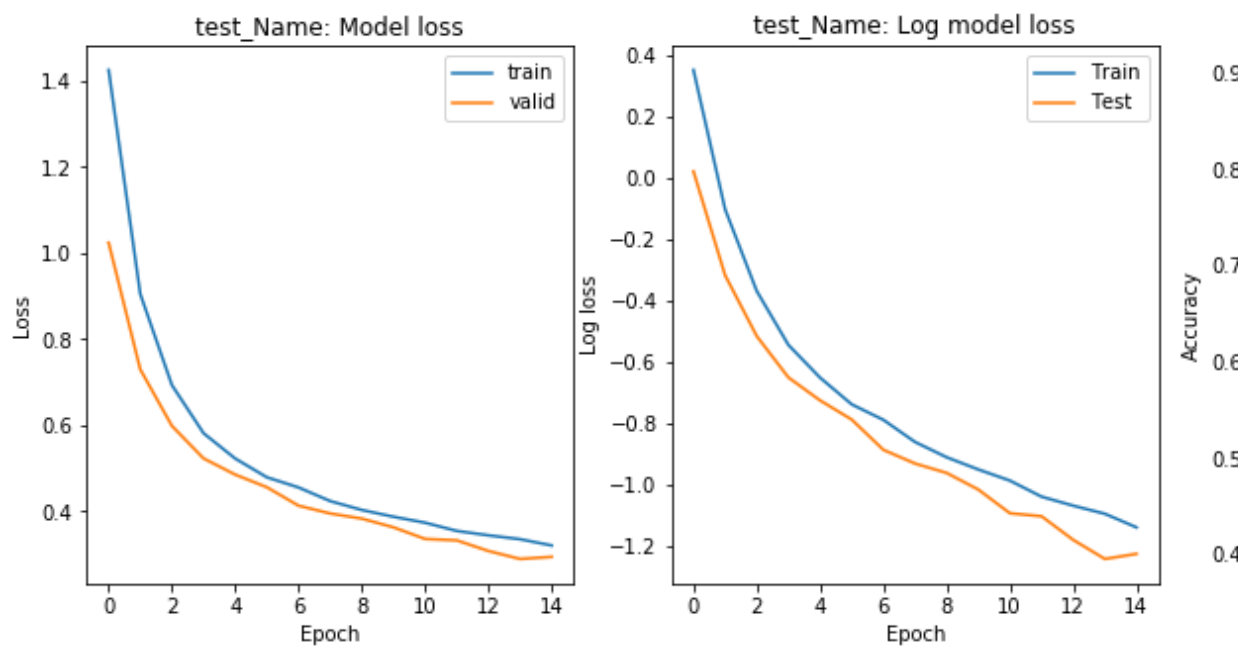
accuracy: 0.8958 - val\_loss: 0.2942 - val\_accuracy: 0.8980

In [167... *## Plot results*


```
plot_history(history, "test_Name")
```

```
# plot confusion matrix
```

```
cnn_helper.plot_confusion_matrix_from_generator(cnn_model, valid_ge
```



Accuracy: 0.8979428266096714





Q: Is there such a thing as too much regularisation?

## Visualise your best CNN

Use the code below to visualise some of the weights you have trained. Hint: Weights are present in convolutional filters and dense layers, nowhere else.

Visualize both one layer with filters, and the outputlayer

```
In [1... # Pick the layer
print(cnn_model.layers)
cw1 = np.array(cnn_model.layers[1].get_weights()) ## Pick the layer w
print(cw1.shape) # 2 weight, 1 weight, 1 bias
print(cw1[0].shape) # Weights
print(cw1[1].shape) # Biases
matrix = cw1[0]


[<tensorflow.python.keras.engine.input_layer.InputLayer object at
0x7f40a427fb70>, <tensorflow.python.keras.layers.convolutional.Conv2D
object at 0x7f40a427f3c8>,
<tensorflow.python.keras.layers.pooling.MaxPooling2D object at
0x7f40a427fa58>, <tensorflow.python.keras.layers.core.Dropout object at
0x7f40a43983c8>,
<tensorflow.python.keras.layers.normalization_v2.BatchNormalization object
at 0x7f40a4398828>, <tensorflow.python.keras.layers.convolutional.Conv2D
object at 0x7f40a4398358>,
<tensorflow.python.keras.layers.pooling.MaxPooling2D object at
0x7f40a46f60f0>, <tensorflow.python.keras.layers.core.Flatten object at
0x7f407c793d30>, <tensorflow.python.keras.layers.core.Dense object at
0x7f40540d1240>]
(2,)
(3, 3, 1, 5)
(5,)

In [169... # Plot your filters
figure, ax = plt.subplots(2, 3, figsize=(14, 10))
figure.suptitle("Weights visualized", fontsize=20)
axes = ax.ravel()

for i in range(0, 5): # Range should be 0 - the number of filters y
    image = matrix[:, :, :, i:i+1]
    image = np.reshape(image, (3, 3)) ## Reshape to the size of you
    axes[i].set_title("Filter" + str(i+1), fontsize=14)
    axes[i].imshow(image)
    axes[i].set_axis_off()

plt.subplots_adjust(wspace=0.05, hspace=0.05)
plt.show()
plt.close()
```

## Weights visualized



## Using existing models

One great thing to do when making a CNN model is to use an architecture that has worked for similar cases. I happen to know that the existing CNN model VGG **1 6** is a good model for these types of images, try that one next.

There are many way of visualising neural networks, see <https://datascience.stackexchange.com/questions/12851/how-do-you-visualize-neural-network-architectures>, but here is one made by Christophe Avenel

## VGG 1 6

```
In [101]: vgg_model = keras.applications.VGG16(  
            include_top=False,  
            weights=None,  
            input_shape=(32, 32, 1),  
            pooling=None,  
        )  
  
In [102]: # add new classifier layers  
flat1 = layers.Flatten()(vgg_model.layers[-1].output)  
class1 = layers.Dense(1024, activation='relu')(flat1)  
output = layers.Dense(5, activation='softmax')(class1)  
  
In [103]: vgg_model = keras.Model(inputs=vgg_model.inputs, outputs=output)  
print (vgg_model.summary())
```

Model: "model"

Layer (type)	Output Shape	Param #
input_31 (InputLayer)	[(None, 32, 32, 1)]	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	640
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160
block4_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block4_conv3 (Conv2D)	(None, 4, 4, 512)	2359808
block4_pool (MaxPooling2D)	(None, 2, 2, 512)	0
block5_conv1 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv2 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv3 (Conv2D)	(None, 2, 2, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0
flatten_17 (Flatten)	(None, 512)	0
dense_17 (Dense)	(None, 1024)	525312
dense_18 (Dense)	(None, 5)	5125
Total params: 15,243,973		
Trainable params: 15,243,973		
Non-trainable params: 0		
None		

Q: How many parameters does this model have?

```
In ... ## Compile the model
```

```
vgg_model.compile(optimizer=keras.optimizers.Adam(learning_rate=0.0001
```

## Q: Why do we need a new classification layers?

### Optional hints for question above

1. What is the original network classifying?
2. What do we want to classify?

### Optional hints for The hint, if you need it

1. So how do we remove the previous classification and make the new one? Just like the code above naturally! A flattening layer is almost always followed by a dense layer or two to expand the model, and then a final classification layer.

```
In [105]: ## Actually train model
          epochs = 10
          history = vgg_model.fit_generator(generator=train_generator,
                                             steps_per_epoch= train_steps,
                                             validation_data= valid_generator,
                                             validation_steps= validation_steps,
                                             epochs= epochs
                                             )
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...
to
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...
to
['...']
```

Train for 1871 steps, validate for 467 steps

Epoch 1/10

1871/1871 [=====] - 34s 18ms/step - loss: 0.7961  
- accuracy: 0.6388 - val\_loss: 0.3651 - val\_accuracy: 0.8959

Epoch 2/10

1871/1871 [=====] - 32s 17ms/step - loss: 0.2542  
- accuracy: 0.9170 - val\_loss: 0.2209 - val\_accuracy: 0.9221

Epoch 3/10

1871/1871 [=====] - 32s 17ms/step - loss: 0.2022  
- accuracy: 0.9369 - val\_loss: 0.1753 - val\_accuracy: 0.9425

Epoch 4/10

1871/1871 [=====] - 32s 17ms/step - loss: 0.1784  
- accuracy: 0.9429 - val\_loss: 0.1382 - val\_accuracy: 0.9553

Epoch 5/10

1871/1871 [=====] - 33s 18ms/step - loss: 0.1654  
- accuracy: 0.9486 - val\_loss: 0.1406 - val\_accuracy: 0.9569

Epoch 6/10

1871/1871 [=====] - 31s 17ms/step - loss: 0.1540  
- accuracy: 0.9502 - val\_loss: 0.1672 - val\_accuracy: 0.9499

Epoch 7/10

1871/1871 [=====] - 32s 17ms/step - loss: 0.1484  
- accuracy: 0.9530 - val\_loss: 0.1467 - val\_accuracy: 0.9561

Epoch 8/10

1871/1871 [=====] - 32s 17ms/step - loss: 0.1435  
- accuracy: 0.9544 - val\_loss: 0.1481 - val\_accuracy: 0.9497

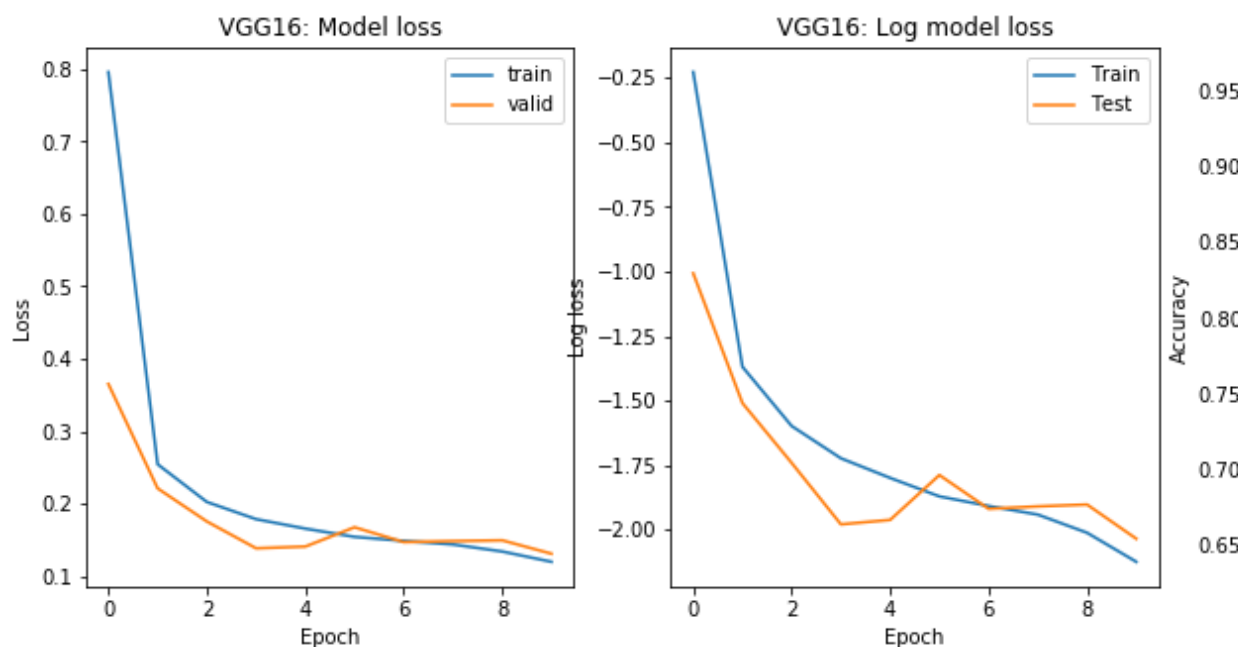
Epoch 9/10

1871/1871 [=====] - 39s 21ms/step - loss: 0.1337  
- accuracy: 0.9568 - val\_loss: 0.1492 - val\_accuracy: 0.9526

Epoch 10/10

1871/1871 [=====] - 33s 17ms/step - loss: 0.1195  
- accuracy: 0.9625 - val\_loss: 0.1307 - val\_accuracy: 0.9599


```
In [106]: ## Plot results
plot_history(history, "VGG16")
```



<Figure size 432x288 with 0 Axes>

```
In [107]: # plot confusion matrix
cnn_helper.plot_confusion_matrix_from_generator(vgg_model, valid_gen)
```

Accuracy: 0.9599251936948971



Q: What is your worst performing class in this classifier? Is it the same as in the other ones?

Q: How many layers with **1 0** filters of size **3 \* 3** would you have to add to the first CNN model we designed to achieve the same number of parameters?

## Try some more models.

Try other optimizers, learning rates, batch sizes or number of epochs. Which would you like to try first and why? Show at least **4** models

```
In [... ## Set up the model architecture
##
# Model try 1
# Accuracy validation: 0.9433
# Comment: Gives good result but the model seems over fitted so for ne
##

cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same', act
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

## Define the model
cnn_model_final = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs,

## Compile the model
cnn_model_final.compile(optimizer=keras.optimizers.Adam(), loss='categ
cnn_model_final.summary()
```

Model: "final\_model"

Layer (type)	Output Shape	Param #
input_8 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_14 (Conv2D)	(None, 32, 32, 5)	50
max_pooling2d_14 (MaxPooling)	(None, 16, 16, 5)	0
conv2d_15 (Conv2D)	(None, 16, 16, 5)	230
max_pooling2d_15 (MaxPooling)	(None, 8, 8, 5)	0
flatten_7 (Flatten)	(None, 320)	0
dense_7 (Dense)	(None, 5)	1605
Total params: 1,885		
Trainable params: 1,885		
Non-trainable params: 0		

```
In [34]: ## Actually train model
         epochs = 10
         history = cnn_model_final.fit_generator(generator=train_generator,
                                                steps_per_epoch= train_steps,
                                                validation_data= valid_generator,
                                                validation_steps= validation_steps,
                                                epochs= epochs,
                                                )
```



WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

Train for 1871 steps, validate for 467 steps

Epoch 1/10

1871/1871 [=====] - 24s 13ms/step - loss: 0.4846

- accuracy: 0.8270 - val\_loss: 0.2676 - val\_accuracy: 0.9141

Epoch 2/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.2395 -

accuracy: 0.9250 - val\_loss: 0.2117 - val\_accuracy: 0.9301

Epoch 3/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.2047 -

accuracy: 0.9357 - val\_loss: 0.2145 - val\_accuracy: 0.9293

Epoch 4/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.1902 -

accuracy: 0.9393 - val\_loss: 0.2147 - val\_accuracy: 0.9296

Epoch 5/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.1801 -

accuracy: 0.9423 - val\_loss: 0.2119 - val\_accuracy: 0.9293

Epoch 6/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.1743 -

accuracy: 0.9430 - val\_loss: 0.1740 - val\_accuracy: 0.9435

Epoch 7/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.1676 -

accuracy: 0.9451 - val\_loss: 0.1884 - val\_accuracy: 0.9392

Epoch 8/10

1871/1871 [=====] - 16s 8ms/step - loss: 0.1651 -

accuracy: 0.9464 - val\_loss: 0.2087 - val\_accuracy: 0.9296

Epoch 9/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.1555 -

accuracy: 0.9507 - val\_loss: 0.1943 - val\_accuracy: 0.9382

Epoch 10/10

1871/1871 [=====] - 16s 8ms/step - loss: 0.1564 -


accuracy: 0.9495 - val\_loss: 0.1887 - val\_accuracy: 0.9433

In [35... *## Plot results*


```
plot_history(history, "final_model")
```

```
# plot confusion matrix
```

```
cnn_helper.plot_confusion_matrix_from_generator(cnn_model_final, val
```



Accuracy: 0.943093775046754



```
In [... ## Set up the model architecture
##
# Model try 2
# Added drop out layer of 0.1 rate.
# Accuracy validation: 0.9384
# Comment: A bit worse accuracy, but this model does not over-fit the
##

cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same', act
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Dropout(.1)(x)
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

## Define the model
cnn_model_final = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs,

## Compile the model
cnn_model_final.compile(optimizer=keras.optimizers.Adam(), loss='categ
cnn_model_final.summary()
```

Model: "final\_model"

Layer (type)	Output Shape	Param #
=====		
input_15 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_28 (Conv2D)	(None, 32, 32, 5)	50
max_pooling2d_28 (MaxPooling	(None, 16, 16, 5)	0
dropout_12 (Dropout)	(None, 16, 16, 5)	0
conv2d_29 (Conv2D)	(None, 16, 16, 5)	230
max_pooling2d_29 (MaxPooling	(None, 8, 8, 5)	0
flatten_14 (Flatten)	(None, 320)	0
dense_14 (Dense)	(None, 5)	1605
=====		
Total params: 1,885		
Trainable params: 1,885		
Non-trainable params: 0		

```
In [55]: ## Actually train model
epochs = 10
history = cnn_model_final.fit_generator(generator=train_generator,
                                         steps_per_epoch= train_steps,
                                         validation_data= valid_generator,
                                         validation_steps= validation_steps,
                                         epochs= epochs,
                                         )
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to  
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to  
['...']
```

Train for 1871 steps, validate for 467 steps

Epoch 1/10

1871/1871 [=====] - 19s 10ms/step - loss: 0.5587  
- accuracy: 0.7994 - val\_loss: 0.3209 - val\_accuracy: 0.8953

Epoch 2/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.3008 -  
accuracy: 0.9033 - val\_loss: 0.2277 - val\_accuracy: 0.9312

Epoch 3/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.2396 -  
accuracy: 0.9228 - val\_loss: 0.2196 - val\_accuracy: 0.9264

Epoch 4/10

1871/1871 [=====] - 16s 8ms/step - loss: 0.2188 -  
accuracy: 0.9309 - val\_loss: 0.2140 - val\_accuracy: 0.9248

Epoch 5/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.2017 -  
accuracy: 0.9352 - val\_loss: 0.1867 - val\_accuracy: 0.9374

Epoch 6/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.1984 -  
accuracy: 0.9382 - val\_loss: 0.1766 - val\_accuracy: 0.9408

Epoch 7/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.1897 -  
accuracy: 0.9379 - val\_loss: 0.1747 - val\_accuracy: 0.9376

Epoch 8/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.1827 -  
accuracy: 0.9410 - val\_loss: 0.1914 - val\_accuracy: 0.9382

Epoch 9/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.1793 -  
accuracy: 0.9401 - val\_loss: 0.1871 - val\_accuracy: 0.9347

Epoch 10/10


1871/1871 [=====] - 17s 9ms/step - loss: 0.1750 -  
accuracy: 0.9425 - val\_loss: 0.1848 - val\_accuracy: 0.9384

In [56... *## Plot results*


```
plot_history(history, "final_model")
```

```
# plot confusion matrix
```

```
cnn_helper.plot_confusion_matrix_from_generator(cnn_model_final, val
```



Accuracy: 0.9382847982901416



```
In [... ## Set up the model architecture
##
# Model try 3
# Changed the Conv2D kernel to see if there is any change in performen
# Accuracy validation: 0.9066
# Comment: Did nit improve the performance, conclude that kernel size
##

cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(4, 4), strides=2,padding='same', act
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Dropout(.1)(x)
x = layers.Conv2D(5, kernel_size=(4, 4), strides=2,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

## Define the model
cnn_model_final = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs,

## Compile the model
cnn_model_final.compile(optimizer=keras.optimizers.Adam(), loss='categ
cnn_model_final.summary()
```

Model: "final\_model"

Layer (type)	Output Shape	Param #
=====		
input_16 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_30 (Conv2D)	(None, 16, 16, 5)	85
max_pooling2d_30 (MaxPooling)	(None, 8, 8, 5)	0
dropout_13 (Dropout)	(None, 8, 8, 5)	0
conv2d_31 (Conv2D)	(None, 4, 4, 5)	405
max_pooling2d_31 (MaxPooling)	(None, 2, 2, 5)	0
flatten_15 (Flatten)	(None, 20)	0
dense_15 (Dense)	(None, 5)	105
=====		
Total params: 595		
Trainable params: 595		
Non-trainable params: 0		

```
In [58]: ## Actually train model
epochs = 10
history = cnn_model_final.fit_generator(generator=train_generator,
                                         steps_per_epoch= train_steps,
                                         validation_data= valid_generator,
                                         validation_steps= validation_steps,
                                         epochs= epochs,
                                         )
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

Train for 1871 steps, validate for 467 steps

Epoch 1/10

1871/1871 [=====] - 23s 12ms/step - loss: 0.9322

- accuracy: 0.6289 - val\_loss: 0.5939 - val\_accuracy: 0.7968

Epoch 2/10

1871/1871 [=====] - 14s 8ms/step - loss: 0.5645 -

accuracy: 0.7974 - val\_loss: 0.4470 - val\_accuracy: 0.8539

Epoch 3/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.4625 -

accuracy: 0.8429 - val\_loss: 0.4179 - val\_accuracy: 0.8643

Epoch 4/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.4213 -

accuracy: 0.8582 - val\_loss: 0.3683 - val\_accuracy: 0.8790

Epoch 5/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.4028 -

accuracy: 0.8693 - val\_loss: 0.3398 - val\_accuracy: 0.8967

Epoch 6/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.3769 -

accuracy: 0.8771 - val\_loss: 0.3328 - val\_accuracy: 0.8969

Epoch 7/10

1871/1871 [=====] - 14s 8ms/step - loss: 0.3717 -

accuracy: 0.8769 - val\_loss: 0.3394 - val\_accuracy: 0.8895

Epoch 8/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.3643 -

accuracy: 0.8807 - val\_loss: 0.3244 - val\_accuracy: 0.8972

Epoch 9/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.3539 -

accuracy: 0.8850 - val\_loss: 0.3001 - val\_accuracy: 0.9066

Epoch 10/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.3535 -


accuracy: 0.8854 - val\_loss: 0.2966 - val\_accuracy: 0.9066

In [61]: *## Plot results*


```
plot_history(history, "final_model")
```

```
# plot confusion matrix
```

```
cnn_helper.plot_confusion_matrix_from_generator(cnn_model_final, val
```



Accuracy: 0.2121293080416778





```

In [... ## Set up the model architecture
##
# Model try 4
# Try another optimizer instead of Adam try SGD
# Accuracy validation: 0.9344
# Comment: Not a whole lot worse than model 2, but the difference is e
##

cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same', act
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Dropout(.1)(x)
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

## Define the model
cnn_model_final = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs,

## Compile the model
cnn_model_final.compile(optimizer=keras.optimizers.SGD(), loss='catego
cnn_model_final.summary()

```

Model: "final\_model"

Layer (type)	Output Shape	Param #
=====		
input_18 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_34 (Conv2D)	(None, 32, 32, 5)	50
max_pooling2d_34 (MaxPooling	(None, 16, 16, 5)	0
dropout_15 (Dropout)	(None, 16, 16, 5)	0
conv2d_35 (Conv2D)	(None, 16, 16, 5)	230
max_pooling2d_35 (MaxPooling	(None, 8, 8, 5)	0
flatten_17 (Flatten)	(None, 320)	0
dense_17 (Dense)	(None, 5)	1605
=====		
Total params: 1,885		
Trainable params: 1,885		
Non-trainable params: 0		

```

In [63]: ## Actually train model
epochs = 10
history = cnn_model_final.fit_generator(generator=train_generator,
                                         steps_per_epoch= train_steps,
                                         validation_data= valid_generator,
                                         validation_steps= validation_steps,
                                         epochs= epochs,
                                         )

```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

Train for 1871 steps, validate for 467 steps

Epoch 1/10

1871/1871 [=====] - 23s 12ms/step - loss: 0.7054

- accuracy: 0.7291 - val\_loss: 0.4113 - val\_accuracy: 0.8627

Epoch 2/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.3928 -

accuracy: 0.8626 - val\_loss: 0.2909 - val\_accuracy: 0.9071

Epoch 3/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.3340 -

accuracy: 0.8844 - val\_loss: 0.3114 - val\_accuracy: 0.8916

Epoch 4/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.2856 -

accuracy: 0.9018 - val\_loss: 0.2333 - val\_accuracy: 0.9240

Epoch 5/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.2620 -

accuracy: 0.9122 - val\_loss: 0.2339 - val\_accuracy: 0.9251

Epoch 6/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.2484 -

accuracy: 0.9171 - val\_loss: 0.2100 - val\_accuracy: 0.9248

Epoch 7/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.2321 -

accuracy: 0.9202 - val\_loss: 0.2411 - val\_accuracy: 0.9234

Epoch 8/10

1871/1871 [=====] - 16s 9ms/step - loss: 0.2269 -

accuracy: 0.9227 - val\_loss: 0.2485 - val\_accuracy: 0.9154

Epoch 9/10

1871/1871 [=====] - 15s 8ms/step - loss: 0.2170 -

accuracy: 0.9279 - val\_loss: 0.1825 - val\_accuracy: 0.9427

Epoch 10/10

1871/1871 [=====] - 14s 8ms/step - loss: 0.2170 -

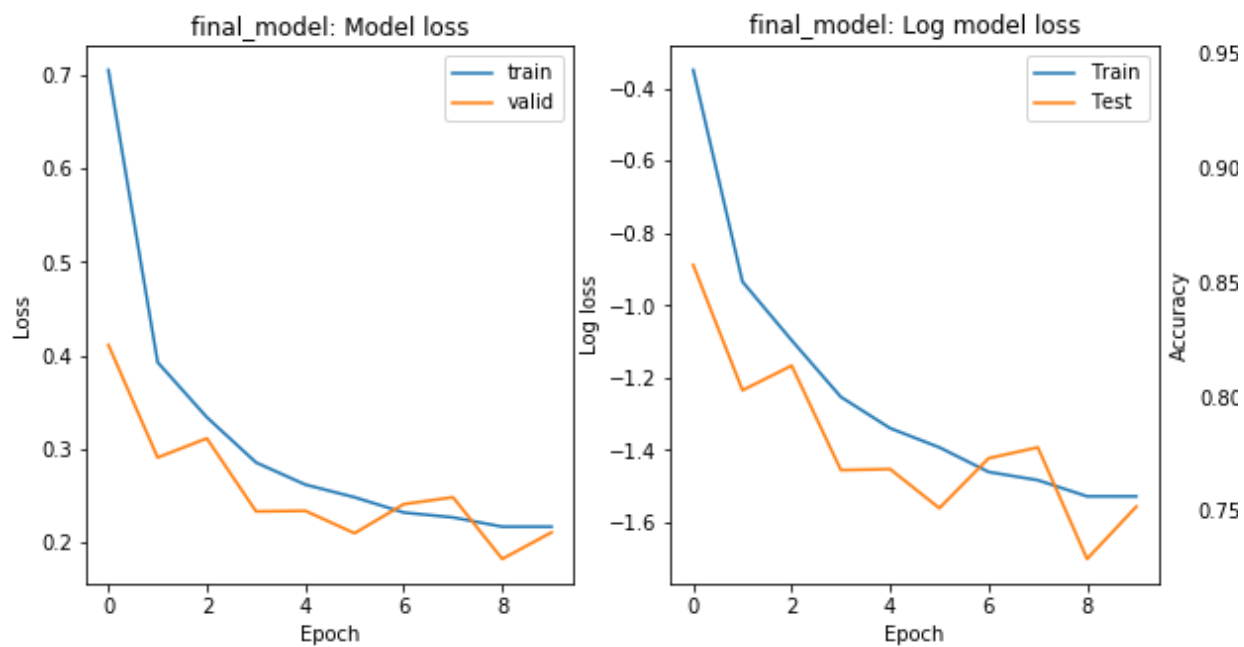
accuracy: 0.9271 - val\_loss: 0.2109 - val\_accuracy: 0.9344

In [64... *## Plot results*

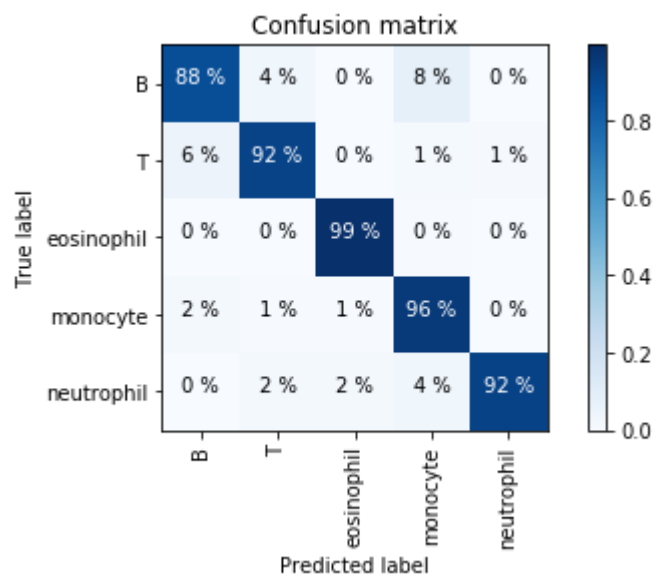
```
plot_history(history, "final_model")
```

```
# plot confusion matrix
```

```
cnn_helper.plot_confusion_matrix_from_generator(cnn_model_final, val
```



Accuracy: 0.9342773176596313



# Finally test your best model

```
In [... ## Set up the model architecture
##
# Took the best model and increased the number of epochs
# Accuracy validation: 0.9406
##

cnn_inputs = keras.Input(shape=(32,32,1))
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same', act
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Dropout(.2)(x)
x = layers.Conv2D(5, kernel_size=(3, 3), strides=1,padding='same')(x)
x = layers.MaxPooling2D(pool_size=(2, 2))(x)
x = layers.Flatten()(x)
cnn_outputs = layers.Dense(5, activation='softmax')(x)

## Define the model
cnn_model_final = keras.Model(inputs=cnn_inputs, outputs=cnn_outputs,

## Compile the model
cnn_model_final.compile(optimizer=keras.optimizers.Adam(), loss='categ
cnn_model_final.summary()
```

Model: "final\_model"

Layer (type)	Output Shape	Param #
=====		
input_20 (InputLayer)	[(None, 32, 32, 1)]	0
conv2d_38 (Conv2D)	(None, 32, 32, 5)	50
max_pooling2d_38 (MaxPooling)	(None, 16, 16, 5)	0
dropout_17 (Dropout)	(None, 16, 16, 5)	0
conv2d_39 (Conv2D)	(None, 16, 16, 5)	230
max_pooling2d_39 (MaxPooling)	(None, 8, 8, 5)	0
flatten_19 (Flatten)	(None, 320)	0
dense_19 (Dense)	(None, 5)	1605
=====		
Total params: 1,885		
Trainable params: 1,885		
Non-trainable params: 0		

```
In [69]: ## Actually train model
epochs = 10
history = cnn_model_final.fit_generator(generator=train_generator,
                                         steps_per_epoch= train_steps,
                                         validation_data= valid_generator,
                                         validation_steps= validation_steps,
                                         epochs= epochs,
                                         )
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

WARNING:tensorflow:sample\_weight modes were coerced from

```
...  
to
```

```
['...']
```

Train for 1871 steps, validate for 467 steps

Epoch 1/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.6162 - accuracy: 0.7646 - val\_loss: 0.3613 - val\_accuracy: 0.8755

Epoch 2/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.3207 - accuracy: 0.8935 - val\_loss: 0.2413 - val\_accuracy: 0.9200

Epoch 3/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.2597 - accuracy: 0.9133 - val\_loss: 0.2039 - val\_accuracy: 0.9358

Epoch 4/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.2381 - accuracy: 0.9204 - val\_loss: 0.2237 - val\_accuracy: 0.9221

Epoch 5/10

1871/1871 [=====] - 17s 9ms/step - loss: 0.2165 - accuracy: 0.9292 - val\_loss: 0.1929 - val\_accuracy: 0.9363

Epoch 6/10

1871/1871 [=====] - 18s 10ms/step - loss: 0.2129 - accuracy: 0.9290 - val\_loss: 0.2168 - val\_accuracy: 0.9205

Epoch 7/10

1871/1871 [=====] - 18s 9ms/step - loss: 0.2042 - accuracy: 0.9317 - val\_loss: 0.2269 - val\_accuracy: 0.9256

Epoch 8/10

1871/1871 [=====] - 18s 10ms/step - loss: 0.2010 - accuracy: 0.9331 - val\_loss: 0.1734 - val\_accuracy: 0.9435

Epoch 9/10

1871/1871 [=====] - 19s 10ms/step - loss: 0.1959 - accuracy: 0.9346 - val\_loss: 0.1663 - val\_accuracy: 0.9451

Epoch 10/10


1871/1871 [=====] - 18s 10ms/step - loss: 0.1934 - accuracy: 0.9361 - val\_loss: 0.1675 - val\_accuracy: 0.9438

In [ ... *## Plot results*

```
plot_history(history, "final_model")
```

```
# plot confusion matrix
```

```
cnn_helper.plot_confusion_matrix_from_generator(cnn_model_final, vali
```




```
In [... test_steps=test_generator.n//test_generator.batch_size if test_genera

pred=cnn_model_final.predict_generator(test_generator,
steps=test_steps,
verbose=1)
```

WARNING:tensorflow:From <ipython-input-221-c5d91644e355>:5:  
Model.predict\_generator (from tensorflow.python.keras.engine.training) is  
deprecated and will be removed in a future version.  
Instructions for updating:  
Please use Model.predict, which supports generators.  
259/259 [=====] - 2s 8ms/step

```
In [... cnn_helper.plot_confusion_matrix_from_generator(cnn_model_final, test_

Accuracy: 0.9379210779595765
```



# ANN

Make a neural network without any convolutions that achieves atleast 90% on the validation test. It will be possible with the techniques you have used above.

```
In [94]: ## Set up the model architecture
inputs = keras.Input(shape= (32,32,1))
#Extend your model here (atleast)
# 32*32=1024
x = layers.Flatten()(inputs)
x = layers.Dense(70, activation=tf.nn.relu)(x)
x = layers.Dropout(.5)(x)
x = layers.Dense(20, activation=tf.nn.relu)(x)
#x = layers.Dropout(.2)(x)
outputs = layers.Dense(5, activation="softmax")(x)
```

```
In [95... ## Define the model
ann_model = keras.Model(inputs=inputs, outputs=outputs, name="ann_M
```

```
In [... ## Compile the model
ann_model.compile(optimizer=keras.optimizers.Adam(), loss='categorical_crossentropy')
ann_model.summary()
```

Model: "ann\_Model"

Layer (type)	Output Shape	Param #
=====		
input_16 (InputLayer)	[(None, 32, 32, 1)]	0
=====		
flatten_15 (Flatten)	(None, 1024)	0
=====		
dense_38 (Dense)	(None, 70)	71750
=====		
dropout_8 (Dropout)	(None, 70)	0
=====		
dense_39 (Dense)	(None, 20)	1420
=====		
dense_40 (Dense)	(None, 5)	105
=====		
Total params: 73,275		
Trainable params: 73,275		
Non-trainable params: 0		

```
In [97]: ## Actually train model
epochs = 25
history = ann_model.fit_generator(generator=train_generator,
                                  steps_per_epoch= train_steps,
                                  validation_data= valid_generator,
                                  validation_steps= validation_steps,
                                  epochs= epochs
                                )
```

WARNING:tensorflow:sample\_weight modes were coerced from

...

to

['...']

WARNING:tensorflow:sample\_weight modes were coerced from

...

to

['...']

Train for 1871 steps, validate for 467 steps

Epoch 1/25

1871/1871 [=====] - 25s 13ms/step - loss: 0.7574

- accuracy: 0.7149 - val\_loss: 0.4422 - val\_accuracy: 0.8295

Epoch 2/25

1871/1871 [=====] - 14s 8ms/step - loss: 0.5127 -

accuracy: 0.8149 - val\_loss: 0.3739 - val\_accuracy: 0.8672

Epoch 3/25

1871/1871 [=====] - 15s 8ms/step - loss: 0.4624 -

accuracy: 0.8310 - val\_loss: 0.3675 - val\_accuracy: 0.8710

Epoch 4/25

1871/1871 [=====] - 15s 8ms/step - loss: 0.4313 -

accuracy: 0.8446 - val\_loss: 0.3414 - val\_accuracy: 0.8812

Epoch 5/25

1871/1871 [=====] - 17s 9ms/step - loss: 0.4030 -

accuracy: 0.8568 - val\_loss: 0.3392 - val\_accuracy: 0.8870

Epoch 6/25

1871/1871 [=====] - 17s 9ms/step - loss: 0.3833 -

accuracy: 0.8616 - val\_loss: 0.3274 - val\_accuracy: 0.8884

Epoch 7/25

1871/1871 [=====] - 15s 8ms/step - loss: 0.3680 -

accuracy: 0.8696 - val\_loss: 0.3303 - val\_accuracy: 0.8897

Epoch 8/25

1871/1871 [=====] - 15s 8ms/step - loss: 0.3544 -

accuracy: 0.8717 - val\_loss: 0.3387 - val\_accuracy: 0.8887

Epoch 9/25

1871/1871 [=====] - 17s 9ms/step - loss: 0.3483 -

accuracy: 0.8735 - val\_loss: 0.3157 - val\_accuracy: 0.8943

Epoch 10/25

1871/1871 [=====] - 17s 9ms/step - loss: 0.3344 -

accuracy: 0.8780 - val\_loss: 0.3331 - val\_accuracy: 0.8878

Epoch 11/25

1871/1871 [=====] - 17s 9ms/step - loss: 0.3283 -

accuracy: 0.8812 - val\_loss: 0.3238 - val\_accuracy: 0.8935

Epoch 12/25

1871/1871 [=====] - 16s 8ms/step - loss: 0.3134 -

accuracy: 0.8895 - val\_loss: 0.3259 - val\_accuracy: 0.8881

Epoch 13/25

1871/1871 [=====] - 17s 9ms/step - loss: 0.2998 -

accuracy: 0.8956 - val\_loss: 0.3128 - val\_accuracy: 0.8943

Epoch 14/25

1871/1871 [=====] - 17s 9ms/step - loss: 0.3054 -

accuracy: 0.8912 - val\_loss: 0.3542 - val\_accuracy: 0.8771

Epoch 15/25

1871/1871 [=====] - 17s 9ms/step - loss: 0.2916 -

accuracy: 0.8979 - val\_loss: 0.3293 - val\_accuracy: 0.8892

Epoch 16/25

1871/1871 [=====] - 15s 8ms/step - loss: 0.2859 -

accuracy: 0.8960 - val\_loss: 0.3289 - val\_accuracy: 0.8878

Epoch 17/25

1871/1871 [=====] - 17s 9ms/step - loss: 0.2833 -

accuracy: 0.8984 - val\_loss: 0.3288 - val\_accuracy: 0.8937

Epoch 18/25

1871/1871 [=====] - 17s 9ms/step - loss: 0.2831 -

accuracy: 0.8985 - val\_loss: 0.3128 - val\_accuracy: 0.8994


Epoch 19/25




1871/1871 [=====] - 16s 9ms/step - loss: 0.2791 -  
accuracy: 0.9015 - val\_loss: 0.3501 - val\_accuracy: 0.8854  
Epoch 20/25  
1871/1871 [=====] - 15s 8ms/step - loss: 0.2715 -  
accuracy: 0.9021 - val\_loss: 0.3238 - val\_accuracy: 0.8932  
Epoch 21/25  
1871/1871 [=====] - 17s 9ms/step - loss: 0.2641 -  
accuracy: 0.9091 - val\_loss: 0.3328 - val\_accuracy: 0.8940  
Epoch 22/25  
1871/1871 [=====] - 17s 9ms/step - loss: 0.2693 -  
accuracy: 0.9058 - val\_loss: 0.3220 - val\_accuracy: 0.8924  
Epoch 23/25  
1871/1871 [=====] - 19s 10ms/step - loss: 0.2628 -  
accuracy: 0.9081 - val\_loss: 0.3266 - val\_accuracy: 0.8935  
Epoch 24/25  
1871/1871 [=====] - 16s 8ms/step - loss: 0.2585 -  
accuracy: 0.9127 - val\_loss: 0.3296 - val\_accuracy: 0.8932  
Epoch 25/25  
1871/1871 [=====] - 16s 9ms/step - loss: 0.2541 -  
accuracy: 0.9102 - val\_loss: 0.3387 - val\_accuracy: 0.8911

```
In [98... ## Plot results
plot_history(history, "test_Name")

# plot confusion matrix
cnn_helper.plot_confusion_matrix_from_generator(ann_model, valid_ge
```



Accuracy: 0.8907293614747529



## Optional

Try using different proportions for training, validation and test. How does this affect your results? Why?

In [ ]: